

linear_regression

February 11, 2019

1 Learn Tensorflow via a linear regression model

In this exercise, you will learn tensorflow in steps. Please use python 3 and tensorflow >=1.4.0.

In [1]: *# A bit of setup*

```
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

from implementations.a_tensorflow import regression_graph
from plot_graph import show_graph

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

1.1 Tensorflow graph

In the first step, we test a tensorflow graph with all constants. Please read `a_tensorflow.regression_graph()`. It calculates the squared loss for the linear regression model.

1.1.1 Question 1 (4 points):

Please run the code in the cell below, comment out the first line, and then run the cell multiple times. Why does the right-most graph get numbers on it nodes with your runs?

1.1.2 Answer:

This is because we comment out the first line so we don't reset the graph. Therefore, everytime we initialize a `regression_graph` object, we actually create a new regression graph in addition to the existing graph. These objects are numbered 1, 2, 3, 4, ...

```
In [2]: #tf.reset_default_graph()
        loss = regression_graph() # the first function in a_tensorflow
        show_graph(loss.graph)
```

<IPython.core.display.HTML object>

Note 1: use `tf.Session()` to run the graph and evaluate the value of a graph node. Note 2: retrieve a graph node by its name.

1.1.3 Question 2 (4 points):

Please uncomment the last two lines and evaluate the loss.

```
In [4]: tf.reset_default_graph()
        loss = regression_graph(print_info=True)
        session = tf.Session()

        # retrieve the variable from the graph
        f = tf.get_default_graph().get_tensor_by_name('score:0')
        print("We retrieve f from the graph and show its value: ", session.run(f))

        print('The loss is ')
        session.run(loss)
```

Msg from the function: we can evaluate any value in the graph with `tf.Session`

Msg from the function: the value of `f = w * x + b` is 2.18

We retrieve `f` from the graph and show its value: 2.18

The loss is

Out[4]: 3.9204001

1.2 Tensorflow graph with vectors

We can get a tensor's shape with its member function `get_shape()`. Note that this function is python function, not a tensorflow operation, so it runs when you build the graph.

We can also use the tensorflow operation `tf.shape()` to get the shape of a tensor. Since it is an operation, it is part of the graph. You need to run the operation to get the actual value.

1.2.1 Question 3 (4 points):

Complete the code in the function `regression_graph_vectorized` so the following cell can run.

```
In [8]: from implementations.a_tensorflow import regression_graph_vectorized
        tf.reset_default_graph()
        loss = regression_graph_vectorized(print_info=True)

        print(tf.Session().run(loss))
```

By `w.get_shape()`, we get the shape of the tensor `w`: (2,)

With `tf.shape(w)`, we get the shape of `w` as a one-element tensor: `Tensor("Shape:0", shape=(1,), dtype=int32)`

The shape of `tf.matmul(x, w)` is: (3, 1)

The shape of `tf.squeeze(tf.matmul(x, w))` is: (3,)

35.343597

1.3 Place holder in Tensorflow graph

A placeholder holds space for the data that will be fed into the graph in the future. It can have shape `None` in one or more dimensions. The size of that dimension will be decided by the data.

1.3.1 Question 4 (4 points):

Please implement the function `regression_graph_with_placeholder` and get the cell below run.

```
In [14]: from implementations.a_tensorflow import regression_graph_with_placeholder
        tf.reset_default_graph()

        # build a graph with place holders
        x = tf.placeholder(shape=[None, 2], dtype=tf.float32, name='feature')
        y = tf.placeholder(shape=[None], dtype=tf.float32, name='label')

        print('The place holder x is:', x) # take a look of the place holder

        loss = regression_graph_with_placeholder(x=x, y=y)

        # need to feed in actual values to x and y to evaluate the variable loss

        # numpy values
        np.random.seed(seed=2019)
        npx = np.random.random_sample([10, 2])
        npy = np.squeeze(npx.dot([[0.3], [0.6]])) + 1.0 + 0.05 * np.random.random_sample([10])

        # this line cannot run because you need to feed in values for place holders. Can you fix it?
        # tf.Session().run(loss)
        l1 = tf.Session().run(loss, feed_dict={x: npx, y: npy})
        print("Loss 1 = ", l1)
```

```

# feed in a different set of data
npx = np.random.random_sample([10, 2])
npy = np.squeeze(npx.dot([[0.7], [0.9]])) + 1.0 + 0.05 * np.random.random_sample([10])

# it get a different loss value
l2 = tf.Session().run(loss, feed_dict={x: npx, y: npy})
print("Loss 2 = ", l2)

```

```

The place holder x is: Tensor("feature:0", shape=(?, 2), dtype=float32)
Loss 1 = 14.350067
Loss 2 = 8.833251

```

1.4 Tensorflow graph with variables

We need to use `tf.Variable` to declare variables that we can optimize later. We often use Variables for model parameters, which are mutable. We seldomly use Variables for data.

1.4.1 Question 5 (4 points):

Please implement the function `regression_graph_with_variable` and get the cell below run.

```

In [16]: from implementations.a_tensorflow import regression_graph_with_variable
         tf.reset_default_graph()

         w = tf.Variable(np.full(shape=[2], fill_value=1.6, dtype=np.float32), name='weight')
         b = tf.Variable(tf.constant(0.9), name='bias')

         loss = regression_graph_with_variable(w=w, b=b)

         session = tf.Session()

         init = tf.global_variables_initializer() # Need to initialize the variables
         session.run(init)

         session.run(loss) # you cannot directly run loss because you need to initialize variables

```

```

Out[16]: 35.343597

```

1.5 Gradient calculation

Tensorflow can calculate gradient with respect to a Variable (but not a constant or a placeholder). Let's use an optimizer to optimize a gradient.

1.5.1 Question 6 (4 points):

Please read the documentation of `compute_gradients` and `apply_gradients`. Can you update `w` and `b` ten times and make the loss smaller?

```

In [20]: # initialize an optimizer
         opt = tf.train.GradientDescentOptimizer(learning_rate=0.01)
         grads_vars = opt.compute_gradients(loss, var_list=[w, b])
         update = opt.apply_gradients(grads_vars)

         print('The current loss is: ', session.run(loss))
         print('The current variable gradients and values are')
         print(session.run(grads_vars))

         for iter in range(10): # can you run multiple steps to minimize the loss?
             session.run(update)

             print('The current loss is: ', session.run(loss))
             # print('The current variable gradients and values are')
             # print(session.run(grads_vars)) # This returns an array corresponding to the gradi

```

```

The current loss is: 0.077929914
The current variable gradients and values are
[(array([ 0.13101178, -0.03523043], dtype=float32), array([1.1932504, 0.8393767], dtype=float32))
The current loss is: 0.07761953
The current loss is: 0.07736203
The current loss is: 0.07713561
The current loss is: 0.07692779
The current loss is: 0.076731496
The current loss is: 0.07654259
The current loss is: 0.07635878
The current loss is: 0.07617866
The current loss is: 0.07600151
The current loss is: 0.07582683

```

1.6 Compose the full version of linear regression

1.6.1 Question 7 (4 points):

Can you use all techniques you have learned above to compose a full version of linear regression?

We use data with one feature only so we can plot the data easily. Assume the feature matrix is `npx` and the label is `npv`.


```

In [37]: # plot the data before fitting the model.
         # numpy values
         N = 100
         npx = np.random.random_sample([N, 1])
         npv = np.squeeze(npx.dot([0.3])) + 1.0 + 0.2 * np.random.random_sample([N])

         plt.plot(npx, npv, 'o')
         plt.ylabel('y')

```

```
plt.xlabel('x')
plt.show()
```



```
In [42]: from implementations.a_tensorflow import regression_graph_full
         tf.reset_default_graph()

         x = tf.placeholder(shape=[None,1], dtype=tf.float32, name='feature')
         y = tf.placeholder(shape=[None], dtype=tf.float32, name='label')

         w = tf.Variable(np.full(shape=[1], fill_value=1.6, dtype=np.float32), name='weight')
         b = tf.Variable(tf.constant(0.9), name='bias')

         loss = regression_graph_full(x=x, y=y, w=w, b=b) # you need to implement this function

         # get an optimizer
         opt = tf.train.GradientDescentOptimizer(learning_rate=0.001)
         # calculate the gradient with compute_gradients
         grads_vars = opt.compute_gradients(loss, var_list=[w, b])
         # get an update operation with apply_gradients
         update = opt.apply_gradients(grads_vars)
```

```

# get a session
session = tf.Session()

# initialize your variables, namely w and b
init = tf.global_variables_initializer() # Need to initialize the variables
session.run(init)

# get feeding dictionary
# npx = np.expand_dims(np_x, axis=1)
# npy = np.expand_dims(np_y, axis=1)

feed_dict={x: npx, y: npy}

print('The current loss is: ', session.run(loss, feed_dict=feed_dict))
print('The current variable gradients and values are')
print(session.run(grads_vars, feed_dict=feed_dict)) # something is missing here?
print(session.run(w))
print(session.run(b))

session.run(update, feed_dict=feed_dict)

for iter in range(1000): # Run 1000 steps to minimize the loss
    # perform gradient steps by running the update operation
    session.run(update, feed_dict=feed_dict)

    if iter % 100 == 1:
        print('Loss values become: ', session.run(loss, feed_dict=feed_dict))
        session.run(grads_vars, feed_dict=feed_dict)

# obtain the trained model
npw = session.run(w)
npb = session.run(b)

```

The current loss is: 32.231148

The current variable gradients and values are

[(array([61.301514], dtype=float32), array([1.6], dtype=float32)), (83.37443, 0.9)]
[1.6]

0.9

Loss values become: 14.168369

Loss values become: 0.9784701

Loss values become: 0.3453107

Loss values become: 0.30256417

Loss values become: 0.29967827

Loss values become: 0.29948324


Loss values become: 0.29947016

```
Loss values become: 0.29946935
Loss values become: 0.29946917
Loss values become: 0.29946917
```

```
In [43]: # Plot the result
```

```
x_line = np.arange(12) / 10.0
y_line = x_line * np.squeeze(npw) + npb
```

```
plt.plot(npx, npy, 'o')
plt.plot(x_line, y_line)
plt.ylabel('y')
plt.ylabel('x')
plt.show()
```



```
In [ ]:
```